



CSCI-GA.2433-001

Database Systems

Lecture 11: MongoDB

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



What Is It?

- A document-oriented database
 - documents encapsulate and encode data (or information) in some standard formats or encodings
- NoSQL database
 - non-adherence to the widely used relational database
 - highly optimized for retrieve and append operations
- uses BSON format
- schema-less
 - No more configuring database columns with types
- No transactions
- No joins

The Basics

- A MongoDB **instance** may have zero or more databases
- A database may have zero or more **collections**.
 - Can be thought of as the relation (table) in DBMS, but with many differences.
- A collection may have zero or more **documents**.
 - Docs in the same collection don't even need to have the same fields
 - Docs are the records in RDBMS
 - Docs can embed other documents
 - Documents are addressed in the database via a unique key
- A document may have one or more **fields**.
- MongoDB **Indexes** is much like their RDBMS counterparts.

The Basics

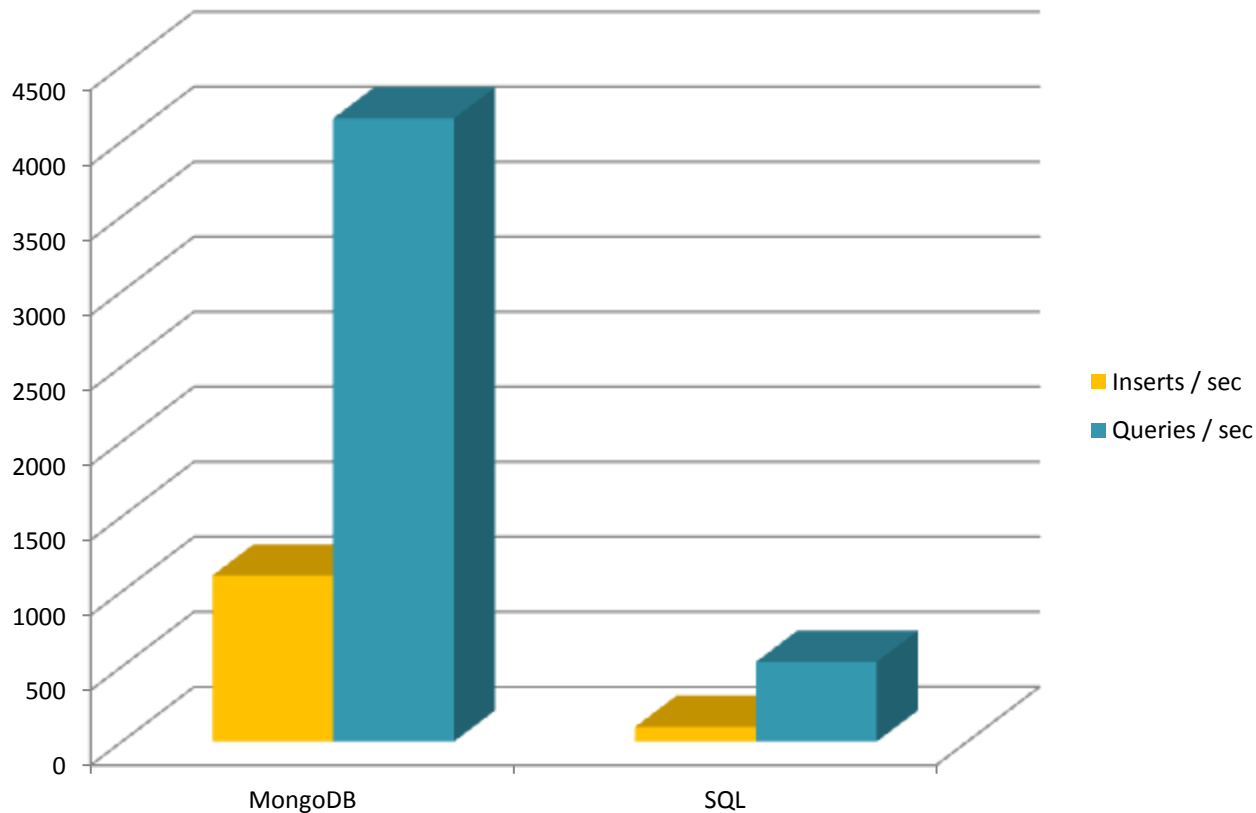
- + Simple queries
- + Makes sense with most web applications
- + Easier and faster integration of data
- Not well suited for heavy and complex transactions systems.

MongoDB Vs Relational DBMS

- Collection vs table
- Document vs row
- Field vs column
- schema-less vs schema-oriented

It Is Fast!

- Anywhere from 2 to 10 times faster than MySQL



Example: Mongo Document

```
user = {  
    name: "Z",  
    occupation: "A scientist",  
    location: "New York"  
}
```

Example: Mongo Document

```
{  
  name: 'Brad Steve',  
  address:  
    {  
      street: 'Oak Terrace',  
      city: 'Denton'  
    }  
}
```


Example: Mongo Collection

```
{ "_id": ObjectId("4efa8d2b7d284dad101e4bc9"),  
  "Last Name": "DUMONT",  
  "First Name": "Jean",  
  "Date of Birth": "01-22-1963" },
```

Obligatory, and
automatically
generated by
MongoDB

```
{ "_id": ObjectId("4efa8d2b7d284dad101e4bc7"),  
  "Last Name": "PELLERIN",  
  "First Name": "Franck",  
  "Date of Birth": "09-19-1983",  
  "Address": "1 chemin des Loges",  
  "City": "VERSAILLES" }
```

Example: A Blog

- A blog post has an author, some text, and many comments
- The comments are unique per post, but one author has many posts
- How would you design this in SQL?

Example: A Blog: Bad Design

- Collections for posts, authors, and comments
- References by manually created ID

```
post = {  
  id: 150,  
  author: 100,  
  text: 'This is a pretty awesome post.',  
  comments: [100, 105, 112]  
}  
author = {  
  id: 100,  
  name: 'Michael Arrington'  
  posts: [150]  
}  
comment = {  
  id: 105,  
  text: 'Whatever this sux.'  
}
```

Example: A Blog: Better Design

- Collection for posts
- Embed comments, author name

```
post = {  
  author: 'Michael Arrington',  
  text: 'This is a pretty awesome post.',  
  comments: [  
    'Whatever this post sux.',  
    'I agree, lame!'  
  ]  
}
```

Why is this one better?

Benefits

- Embedded objects brought back in the same query as parent object
 - Only 1 trip to the DB server required
- Objects in the same collection are generally stored contiguously on disk
 - Spatial locality = faster
- If the document model matches your domain well, it can be much easier to comprehend than nasty joins

Queries in MongoDB

- Query expression objects indicate a pattern to match
 - `db.users.find({last_name: 'Smith'})`
- Several query objects for advanced queries
 - `db.users.find({age: {$gte: 23}})`
 - `db.users.find({age: {$in: [23,25]}})`
- Exact match an entire embedded object
 - `db.users.find({address: {street: 'Oak Terrace', city: 'Denton'}})`
- Dot-notation for a partial match
 - `db.users.find({"address.city": 'Denton'})`

Indexing

- Indexes in MongoDB are similar to indexes in RDBMS.
- MongoDB supports indexes on any field or sub-field contained in documents
- MongoDB defines indexes on a per-collection level.
- All MongoDB indexes use a B-tree data structure.

Building An Application

- Want to build an app where users can check in to a location
- Leave notes or comments about that location
- Requirements
 - Need to store locations (Offices, Restaurants etc)
 - Want to be able to store name, address and tags
 - Maybe User Generated Content, i.e. tips / small notes ?
 - Want to be able to find other locations nearby
 - User should be able to 'check in' to a location
 - Want to be able to generate statistics

Building An Application

loc1, loc2, loc3

Locations

User1, User2

Users

Places

```
location1 = {  
  name: "10gen HQ",  
  address: "17 West 18th Street 8th Floor",  
  city: "New York",  
  zip: "10011",  
  latlong: [40.0,72.0],  
  tags: ["business", "cool place"],  
  tips: [  
    {user:"nosh", time:6/26/2010, tip:"stop by for office hours on  
      Wednesdays from 4-6pm"},  
    {.....},  
  ]  
}
```

Example queries:

- `db.locations.find({latlong:{$near:[40,70]}})`
- `db.locations.find({zip:"10011", tags:"business"})`
- `db.locations.find({zip:"10011"}).limit(10)`

Inserting and updating locations

Initial data load:

```
db.locations.insert(place1)
```

Using update to Add tips:

```
db.locations.update({name:"10gen HQ"},  
  {$push :{tips:  
    {user:"nosh", time:6/26/2010,  
      tip:"stop by for office hours on  
        Wednesdays from 4-6"}}}}}
```

Users

```
user1 = {  
  name: "nosh"  
  email: "nosh@10gen.com",  
  .  
  .  
  .  
  checkins: [{ location: "10gen HQ",  
    ts: 9/20/2010 10:12:00,  
    ...},  
    ...  
  ]  
}
```

Simple Stats

```
db.users.find({'checkins.location': "10gen HQ"})
```

```
db.checkins.find({'checkins.location': "10gen HQ"})  
                .sort({'ts':-1}).limit(10)
```

```
db.checkins.find({'checkins.location': "10gen HQ",  
                ts: {$gt: midnight}}).count()
```

Limitations of MongoDB

- No referential integrity
- High degree of denormalization means updating something in many places instead of one
- Lack of predefined schema is a double-edged sword
 - You must have a model in your app
 - Objects within a collection can be completely inconsistent in their fields

Conclusions

- MongoDB is fast
 - Very little CPU overhead
 - MongoDB is Implemented in C++ for best performance
- Very rapid development, open source
- useful when working with a huge quantity of data when the data's nature does not require a relational model
- used when what really matters is the ability to store and retrieve great quantities of data, not the relationships between the elements.
- Works on many platforms and there are many language drivers

<http://www.mongodb.org/>