



The MongoDB Tutorial

Introduction for MySQL Users

Stephane
Combaudon
April 1st, 2014



PERCONA
LIVE



Agenda

2

- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding



Before we start

3

- No knowledge of MongoDB is required
- Several interesting topics will intentionally be left out
- Any problem, any question? Raise your hand!



Agenda

4

- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding

Typical problems with RDBMS

5

- SQL is difficult to understand
- Changing the schema is difficult with large tables
 - aka “The relational model is not flexible”
- Scaling out and HA are difficult
 - Support for sharding is limited or non-existent
 - Many HA solutions, which one to choose?

MongoDB solutions (1)

6

- “SQL is difficult to understand”
 - Everything is a JSON document in MongoDB
 - No joins but a powerful way to write complex queries (aggregation framework)
 - Queries are easy to read/write
 - SQL: `SELECT * FROM people`
`WHERE name = 'Stephane'`
 - MongoDB: `db.people.find({name: 'Stephane'})`

MongoDB solutions (2)

7

- “The relational model is not flexible”
 - MongoDB is schemaless, no ALTER TABLE!

```
db.people.insert({name: 'Stephane'});
```

```
db.people.insert({name: 'Joe', age: 30});
```

- But be careful, this is *also* allowed

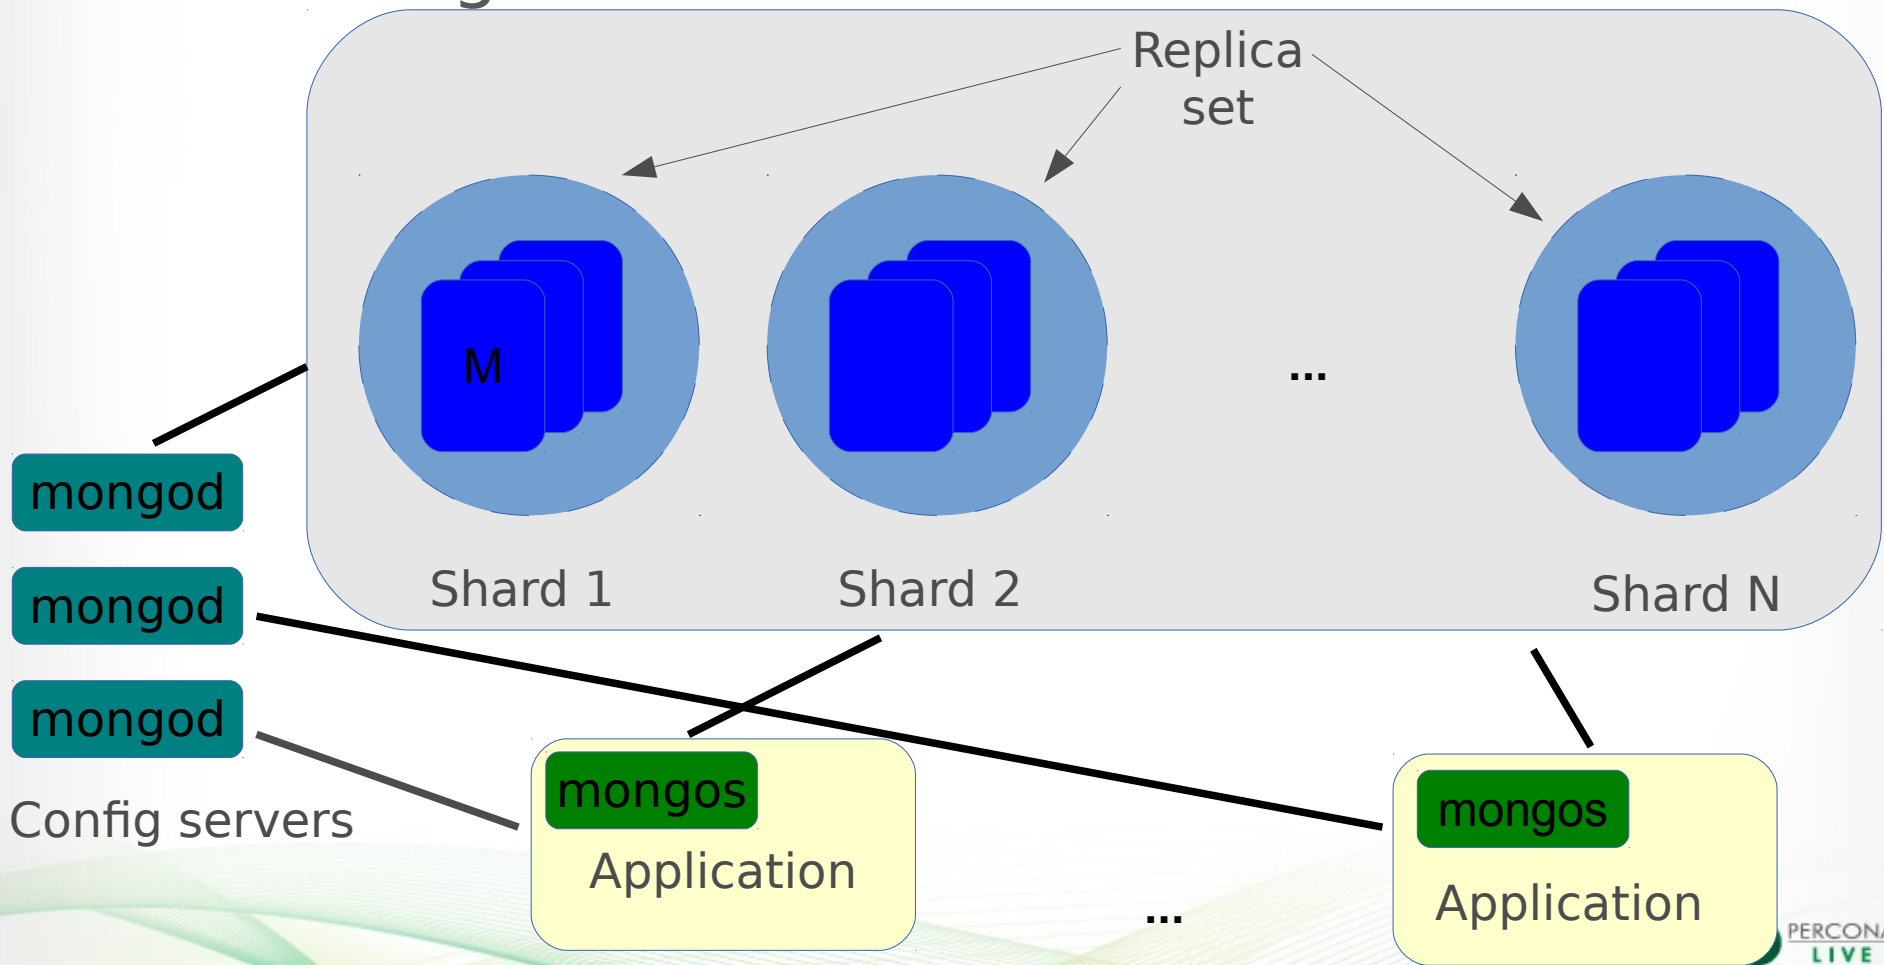
```
db.people.insert({name: 'Stephane'})
```

```
db.people.insert({n: 'Joe'})
```

MongoDB solutions (3)

8

- “Scaling and HA are difficult”



Summary

9

Scalability &
Performance

- Memcached

- MongoDB

- RDBMS

Functionality



Agenda

10

- Introduction
- **Install & First Steps**
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding

MongoDB setup

11

- See `mongod_setup.rst`

Useful terminology

12

Relational	MongoDB
Database	Database
Table	Collection
Row	(JSON) document



JSON

13

- Stands for JavaScript Object Notation
 - But not tied to Javascript
- Open standard for human-readable data interchange
- Lightweight alternative to XML
- Internally, docs are stored in BSON
 - Binary JSON
 - See <http://bsonspec.org>

Invoking the shell

14

\$ mongo (--port=xxx, default 27017)

```
> show dbs
local   0.03125GB
test    0.0625GB
> use test
switched to db test
> show collections
people
system.indexes
> db.people.find()
{ "_id" : ObjectId("523ef7bf8108101415e7d1d1"), "age" : 30, "country" : "XXX", "name" : "Joe"
}
{ "_id" : ObjectId("523ef7ac8108101415e7d1d0"), "age" : 33, "country" : "XXX", "name" : "Steph
ane" }
> db.people.find().pretty()
{
  "_id" : ObjectId("523ef7bf8108101415e7d1d1"),
  "age" : 30,
  "country" : "XXX",
  "name" : "Joe"
}
{
  "_id" : ObjectId("523ef7ac8108101415e7d1d0"),
  "age" : 33,
  "country" : "XXX",
  "name" : "Stephane"
}
```

Inserting data (1)

15

```
use test
```

```
db.people.insert({name:'Stephane',country:'FR'})
```

- This inserts a document
 - In the people collection
 - Collection is in the test database
- Collection is created if it did not exist

Inserting data (2)

16

- The shell embeds a JS interpreter
 - You can also use a variable to build the JSON document step by step

```
> x = {name:'Stephane'}
{ "name" : "Stephane" }
> x.country = 'FR'
FR
> x
{ "name" : "Stephane", "country" : "FR" }
> db.people.insert(x)
> db.people.findOne()
{
  "_id" : ObjectId("526fd69e64ea1df71b82f55b"),
  "name" : "Stephane",
  "country" : "FR"
}
```

- Unique identifier of a doc
- You can set it explicitly

```
db.people.insert({'_id': 'Stephane', country: 'FR'})
```
- If you don't, it will be created for you

```
"_id" : ObjectId("523ef7bf8108101415e7d1d1")
```
- Monotonically increasing on a single node
 - Think auto_increment in MySQL

Structure of a document

18

- Set of key/value pairs
 - `{'key1':'value1','key2':'value2',...}`
- A value can be an array
- A value can be another document



Lab #1

19

- See lab1.rst



Agenda

20

- Introduction
- Install & First Steps
- **CRUD**
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding

Inserting a document

21

- SQL
 - `INSERT INTO t (fn, ln) VALUES ('John', 'Doe')`
- MongoDB
 - `db.t.insert({fn:'John', ln: 'Doe'})`
- No need to specify non-existing fields

Updating a document (1)

22

- SQL

- `UPDATE t SET ln='Smith' WHERE ln='Doe'`

- MongoDB

- `db.t.update({ln:'Doe'},{$set:{ln:'Smith'}})`

- Only 1 doc. is updated by default
 - Specify `{multi:true}` for multi-doc updates

- `db.t.update({ln:'Doe'},{$set:{ln:'Smith'}},{multi:true})`

Updating a document (2)

23

- To add a new field
 - No need to run ALTER TABLE!
 - It is a regular update
 - `db.t.update({ln:'Smith'},{$set:{age:30}})`

Removing documents

24

- SQL

- `DELETE FROM t WHERE fn='John'`

- MongoDB

- `db.t.remove({fn:'John'})`

Dropping a collection

25

- SQL
 - `DROP TABLE t`
- MongoDB
 - `db.t.drop()`

Selecting all documents

26

- SQL
 - `SELECT * FROM t`
- MongoDB
 - `db.t.find()`
- SQL
 - `SELECT id, name FROM t`
- MongoDB
 - `db.t.find({}, {name:1})`

Specifying a “WHERE” clause

27

- SQL

- `SELECT * FROM t WHERE fn='John'`

- MongoDB

- `db.t.find({ln:'John'})`

- SQL

- `SELECT id, age FROM t WHERE fn='John' AND ln='Smith'`

- MongoDB

- `db.t.find({fn:'John',ln:'Smith'},{age:1})`

Sorts, limits

28

- SQL

- `SELECT country FROM t`
`WHERE fn='John'`
`ORDER BY age DESC LIMIT 10`

- MongoDB

- `db.t.find(`
`{fn:'John'},`
`{country:1, _id:0}`
`).sort({age:1}).limit(10)`

Inequalities

29

- SQL

- `SELECT fn,ln FROM t`
`WHERE age > 30`

- MongoDB

- `db.t.find(`
`{age:{$gt:30}},`
`{fn:1, ln:1, _id:0}`
`)`

Conditions on subdocuments

30

```
{  
  _id:....  
  fn:'John', ln:'Doe'  
  address:{  
    country:'US',  
    state:'CA'  
  }  
}
```

```
db.t.find(  
  {address.country:'US'},  
  {fn:1, ln:1, _id:0}  
)
```

Conditions on arrays

31

```
{  
  _id:...  
  hobbies:['music', 'MongoDB', 'Python']  
}
```

- Exact match on array

```
db.t.find({hobbies:['music', 'MongoDB', 'Python']})
```

- Match on array element

```
db.t.find({hobbies:'Python'})
```

Counting

32

- SQL
 - `SELECT COUNT(*) FROM t`
`WHERE country='US'`
- MongoDB
 - `db.t.count({country:'US'})`



Lab #2

33

- See lab2.rst file



Agenda

34

- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding

Aggregation

35

- “Means” GROUP BY, SUM(), ...
- Not possible with the operators we saw previously (exception: count)
- 2 ways to aggregate
 - Map-Reduce # Not covered here
 - Aggregation Framework

Aggregation Framework

36


- Available from MongoDB 2.2
- Easier and faster than Map-Reduce
- Some limitations
 - AF is perfect for simple cases
 - Map-Reduce can be needed for complex situations

10,000 feet overview

37

- Documents are modified through pipelines
- Similar to Unix pipelines

```
grep oo /etc/passwd | sort -rn | awk -F ':' '{print $1,$3,$4}'
```



First example

38

- SQL

```
SELECT fn, count(*) AS total FROM people WHERE fn >= 'M%' GROUP BY fn
```

- MongoDB

```
db.people.aggregate({$match:{fn:{$gte:'M'}}},{$group:{_id:"$fn",total:{$sum:1}}})
```

- You probably want some explanation!

First example, explained

39

```
db.people.aggregate(  
  {$match:XXX},      # Pipeline 1, filtering criteria  
  {$group:XXX}       # Pipeline 2, group by  
)
```

- Filtering condition like with find()

```
$match:{fn:{$gte:'M'}}
```

- Specifying the grouping field

```
$group:{_id:"$fn",...}
```

- Specifying the aggregated fields

```
$group:{..., total:{$sum:1}}
```

Pipeline order matters

40

```
db.people.aggregate(  
  { $match: ... },  
  { $group: ... }  
)
```

vs

```
db.people.aggregate(  
  { $group: ... },  
  { $match: ... }  
)
```

SQL analogy

41

```
SELECT ... FROM people  
WHERE ...  
GROUP BY ...
```

VS

```
SELECT ... FROM people  
GROUP BY ...  
HAVING ...
```

Clean the output

42

- The projecting operator allows renaming keys/values

```
db.people.aggregate(  
  {$match:...},  
  {$group:...},  
  {$project:{_id:0, name:{$toUpper:"$_id"}, total:1}}  
)
```


Final query

43

- SQL

```
SELECT UPPER(fn) AS name, count(*) AS total
FROM people
WHERE fn >= 'M%'
GROUP BY fn
```

- MongoDB

```
db.people.aggregate(
  {$match:{fn:{$gte:'M'}}},
  {$group:{_id:"$fn",total:{$sum:1}}},
  {$project:{_id:0, name:{$toUpper:"$_id"}, total:1}}
)
```

Other operators

44

- **\$sort**
 - `{ $sort: { total: -1 } }`
- **\$limit**
 - `{ $limit: 10 }`
- **\$skip**
 - `{ $skip: 100 }`
- **\$unwind**
 - Splits a document with an array into multiple documents



Lab #3

45

- See lab3.rst



Agenda

46

- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- **Performance Tuning**
- Replication and High Availability
- Sharding

Indexes

47

- MongoDB supports indexes
 - At the collection level
 - Similar to indexes on RDBMS
- Can be used for
 - More efficient filtering
 - More efficient sorting
 - Index-only queries (covering index)

Main types of indexes

48

- Single-field/multiple field index
 - `db.t.ensureIndex({fn:1})`
 - `db.t.ensureIndex({fn:1, age:-1})`
- Unique index
 - All collections have an index on `_id`
 - `db.t.ensureIndex({username:1}, {unique:true})`
- Indexes on arrays, embedded fields, subdocuments are also supported

Compound indexes (1)

49

- Sort order
 - Traversal in either direction

```
db.t.find().sort({country:1,age:-1})
```

- Good indexes

- {country:1, age:-1}
- {country:-1, age:1}

- Bad index

- {country:1, age:1}

Compound indexes (2)

50

- Prefixes
 - Leftmost prefixes are supported

```
db.t.ensureIndex({age:1, country:1})
```

- Good for

- `db.t.find({age:30})`
- `db.t.find({age:30, country:'US'})`

- Not good for

- `db.t.find({country:'US'})`



Lab #4

51

- See lab4.rst



Agenda

52


- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- Sharding



Overview

53

- Replication is asynchronous
- All writes go to the master
- Secondary can accept reads
- A new master is elected if the current master fails
- An arbiter (no data) can be set up for the election process



Setup of a 3 node replica set

55

- Start 3 instances with `--replSet rsname`
- Then on the master

```
rsconf = {  
    _id: "RS",  
    members: [{  
        _id: 0,  
        host: "localhost:30001"  
    }]  
}  
  
rs.initiate(rsconf)  
rs.add("localhost:30002")  
rs.add("localhost:30003")
```

If the master crashes

56

- If a heartbeat does not return within 10s, the master is considered unavailable
- Election of a new master starts then
- You can influence the choice by setting a priority (between 0 and 100)

Setting priorities

57

```
cfg = rs.conf()  
cfg.members[1].priority = 0  
cfg.members[2].priority = 10  
rs.reconfig(cfg)
```

- A node set to have higher priority than the master will be elected the new master

Write concerns (aka w option)

58

- How many nodes should ack a write?
 - $w=1$
 - Primary only (default setting)
 - $w=2$
 - Primary and one secondary
 - $w=\text{majority}$
 - Majority of the nodes

Setting a write concern

59

```
cfg = rs.conf()  
cfg.settings.getLastErrorDefaults =  
{w: "majority"}  
rs.reconfig(cfg)
```

Read preferences

60

- Can be any of
 - primary
 - primaryPreferred
 - secondary
 - secondaryPreferred
 - nearest
 - Or you can use a custom tag



Lab #5

61

- See lab5.rst



Agenda

62

- Introduction
- Install & First Steps
- CRUD
- Aggregation Framework
- Performance Tuning
- Replication and High Availability
- **Sharding**

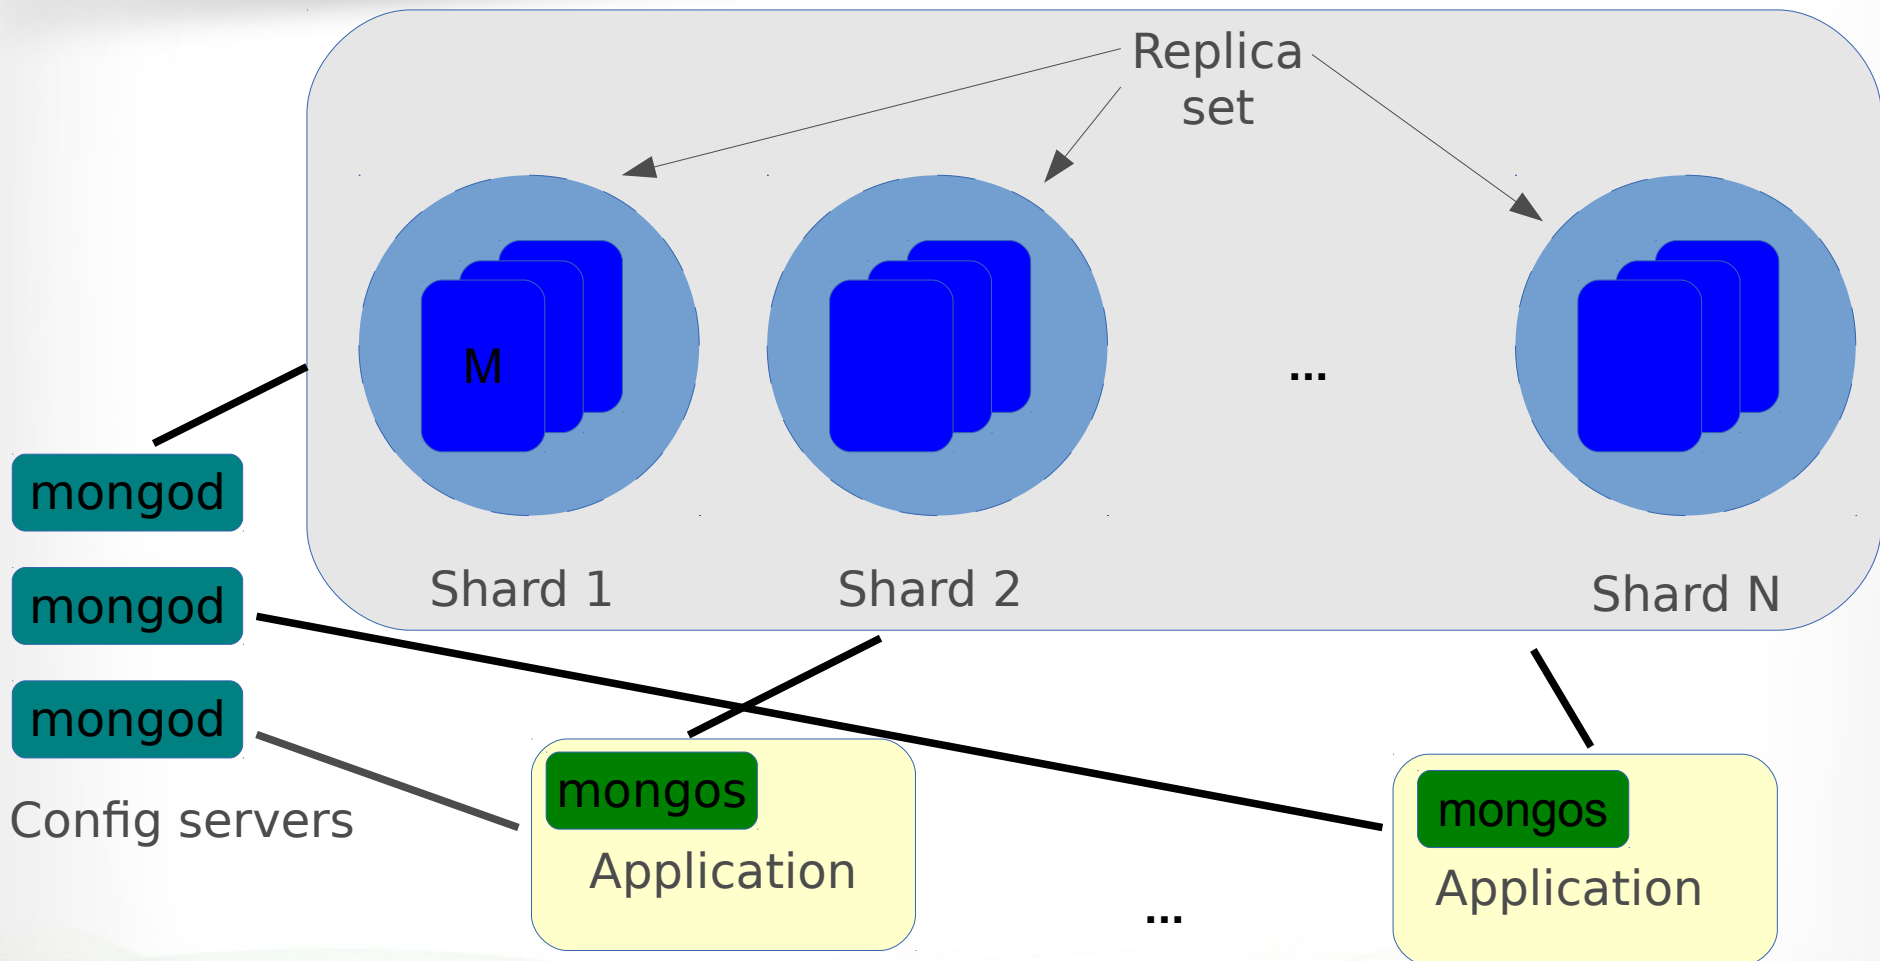
When to shard

63

- When the dataset no longer fits in a single server
- When write activity exceeds the capacity of a single node
- When the working set no longer fits in memory

Architecture of a sharded cluster

64



Setup (1)

65

- Start the config servers
 - `mongod --configsvr --dbpath <path> --port 40001`
 - `mongod --configsvr --dbpath <path> --port 40002`
 - `mongod --configsvr --dbpath <path> --port 40003`
- Start a router
 - `mongos -configdb localhost:40001, localhost:40002, localhost:40003 --port 31000`
- Connect to mongos
 - `mongo --host localhost --port 31000`

Setup (2)

66

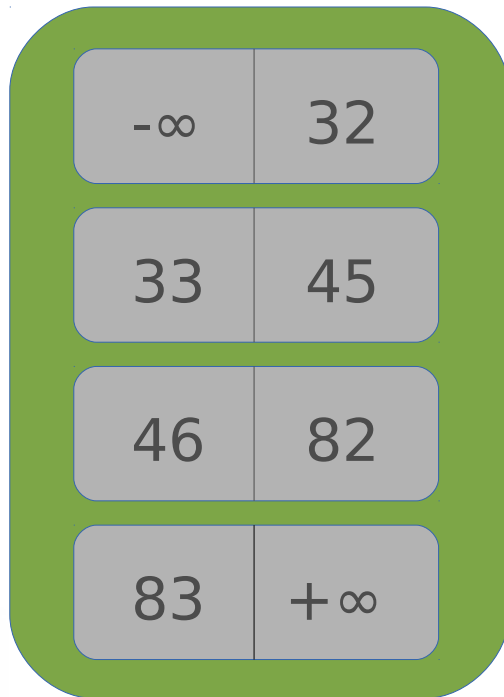
- Add shards
 - `sh.addShard('RS/localhost:30001')`
- Enable sharding for a database
 - `sh.enableSharding('test')`
- Enable sharding for a collection
 - `sh.shardCollection('test.people',{'fn':1})`

Balancing (1)

67

- Data is stored in chunks

```
sh.addShard(shard1)
```



Shard 1

Balancing (2)

68

`sh.addShard(shard2)`



Shard 1



Shard 2

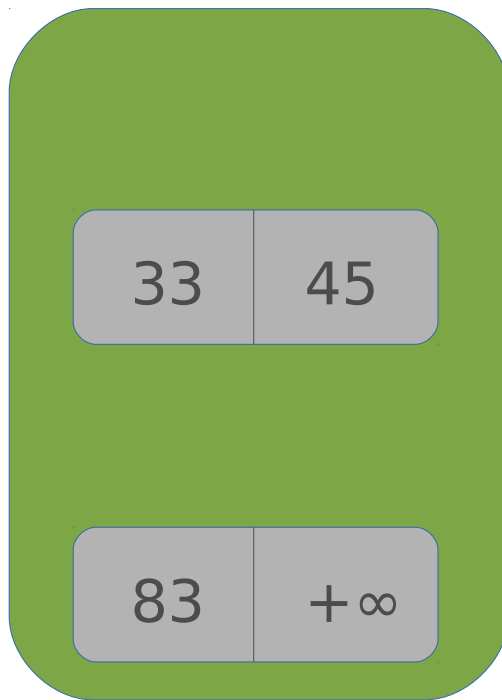
Balancing (3)

69


`sh.addShard(shard3)`



Shard 1



Shard 2



Shard 3



Lab #6

70

- See lab6.rst

Thank you for your attention!

stephane.combaudon@percona.com